

in - input is in file
 out - creek/overwrite file
 fort - append to existing file
 com - command → pipe

[0-2A-Z][0-9]
 ^ - line beginning, \$ - end of line
 \ - start of word, > - end
 * - 0 or more apparitions
 + - 1 or more apparitions
 ? - 0 or 1 apparition
 \ - blank character
 (abc) - group, / - or

grep - g (quiet)
 -c: count the number of matches
 -v: only those not matching
 -m: display matching lines & the line no.
 -r: recursively
 -i: case insensitive
 -x: match exactly the entire line
 -o: only the part that matches

sort - r: reverse
 -u: uniq
 -M: compares lines

sed - f: from a file
 p: print | d: delete
 i: insert (before)
 a: append (after)
 s/regex/string/flag/opt

awk - if, for, while = C
 BEGIN { } END { }
 NF - no. of fields \$1, \$2, ...
 NR -
 \$0 - entire line
 FS - field separator
 length (string)
 split (string, array, sep.)
 index (s1, s2) → string

wc -l: lines
 -w: words -c: characters

Shell /bin/bash
 chmod +x file
 if ... ; then ... ; fi
 \$ var → access the value
 for var in ... ; do ... done

" " - interpreted (spec. chars)
 ' ' - quote a char
 \$ "var1" = \$ "var2"
 test -b, -k, -g, -e, -f, -s, -d, -r, -w, -x
 test -z (string) -n (string)

shell / \$(expr) kex
 command - executes command and replaces it with the result of the execution
 shift - shift arguments to the left
 \$# - argc, \$@ - argv

UNIX - interactive, time-sharing, multi-user, multi-tasking
 KERNEL - nucleus (could use loaded at the PC startup)
 System CALL: Kernel mode to user mode for access to disk, network, video mem.

KERNEL MODE = level 0 protection
 USER MODE = level 3 protection
 levels 1, 2 = services

FILE TYPES: regular, directory, hard/soft links, FIFO, socket
 SYMBOLIC LINKS: ln -s target name - behave as the target file/directory - can reference on another disk

HARD LINKS: ln target file - now i-mode for the same file - file is deleted only when there are no more links to it

ROOT (SUPERUSER): root
 METHODE DE ALLOCARE A MEM
 ① SINGLE TASKING: - one single program, in space another but, we need to share this one

② MULTITASKING:
 A) REAL: a) absolute i) fixed partitions
 + more proc
 + the max size of proc
 + program compiled on a specific partition

ii) partitions relocable: the addresses are stored as OFFSETS inside a partition
 b) dynamic partitions: fragmented memory → memory needs to be compressed (compactata?) (close processes): all of it or as much as needed for one process

→ SOLUTION: paged memory

③ VIRTUAL: a) paged - virtual addresses are generated in the form of (page, offset), each page has a table of pages address: offset + p * dmin - page
 b) segmented: code + data segments - the access to the seg. is controlled by the OS, the seg. can be restricted (swapped from another program)

c) paged + segmented - segments are divided into more virtual pg. the proc. has a table of segments, each of which has a table of pages

LOADING POLICIES
 ① all pages in the beginning
 ② fast access at runtime
 ③ slow startup
 ④ loading pages that are not necessarily needed
 ⑤ each page when needed
 ⑥ fast startup
 ⑦ loading only what is needed
 ⑧ slow access at time

LOCATION PRINCIPLE: the pages neighbouring the page that has just been loaded are more likely to be needed in the future so we load them preemptively

REPLACING POLICIES
 how to choose the victim pg
 FIFO: each pg. has an age, the oldest ones go first
 NRU: each pg. has bits (r, w) that are set to 1 at n/w and the OS is periodically setting them to 0 → classes (0,0), (0,1), (1,0), (1,1)
 first record
 LRU - last recently used
 matrix of N x N, N = no. of pg; when accessing pg. i, line i is populated with 1, columns → 0; pages with the lower 1s are first to go

PIPE pipe (int p[2])
 0 - read, 1 - write

FIFO mbufs (char * name, int mode)
 empty - read blocks until no more data
 full - write blocks until all data is read → spaced
 0 - NOBUF → not 0 if fifo empty
 full
 no process to read/write
 0 - NOBUF → error
 - normal → wait
 - pipes need not be related

STATES OF A PROCESS
 Hold → Ready → Run → Fin
 1 - wait
 2 - wait
 3 - wait
 4 - wait
 5 - wait
 6 - wait

1 - create proc, 2 - get CPU, 3 - release CPU, 4 - move on disk, release CPU, 5 - i/o, wait for resources, 6 - op. from 5 done, free disk space (deallocate memory), 7 - release CPU, free mem, destroy process
 * - alloc. mem.
 * - free mem.

PLACEMENT POLICIES
 how to handle malloc?
 - FIRST FIT: ① fast ② no attempt to control fragmentation
 - BEST FIT: ① economic in terms of large memory areas / chunks ② never creates small areas of free memory → fragmentation
 - WORST FIT: ① leaves behind larger chunk ② never

- BUDDY → always alloc. a chunk of the smallest power of two, greater than the required size; keep lists of free chunks by powers of 2; alloc. the largest available chunk and distribute the remaining memory in pair of 2 chunks
 + speed

VHIT - suspends the exec until a child finishes (any, or a specific one)
 EXEC - replace the whole image with the given command
 EXIT - exits the current proc and goes back to the parent

CAUSES OF DL!
 1. Mutual exclusion
 2. Hold and wait
 3. No preemption
 4. Circular wait

SOLUTION TO DL
 - reset → pick a victim proc → create a backup state
 DETECTION OF DL
 graf: □ proc, ○ res.
 → is held by, -- → is req. cycle → deadlock
 PREVENTION OF DL
 → pick an order for and always lock in the order

ZOMBIE → process that has finished execution but still has an entry in the proc table; wait to avoid or before creating: signal (SIGCHLD, SIG-IGN)
 ORPHAN → still execute parent is dead

SCHEDULING ALGO
 ① FIRST COME, FIRST SERVED
 ② SHORTEST FIRST: we need to know durations
 ③ not always right
 ④ Priority Based
 + starvation
 ⑤ Deadline Scheduling based on the finish time (must finish by)

⑥ Round Robin: time sharing, procs are working simultaneously, in READY, one → R

SEMAPHORE (value with two operations)
 P(s) - wait for proc A
 V(s) -
 if (s < 0)
 state(s) = wait
 q(s) ← A / A wait
 pass ctrl to sys

EXEC - replace the whole image with the given command
 EXIT - exits the current proc and goes back to the parent

PIPE pipe (int p[2])
 0 - read, 1 - write

STATES OF A PROCESS
 Hold → Ready → Run → Fin
 1 - wait
 2 - wait
 3 - wait
 4 - wait
 5 - wait
 6 - wait

PLACEMENT POLICIES
 how to handle malloc?
 - FIRST FIT: ① fast ② no attempt to control fragmentation
 - BEST FIT: ① economic in terms of large memory areas / chunks ② never creates small areas of free memory → fragmentation
 - WORST FIT: ① leaves behind larger chunk ② never

- BUDDY → always alloc. a chunk of the smallest power of two, greater than the required size; keep lists of free chunks by powers of 2; alloc. the largest available chunk and distribute the remaining memory in pair of 2 chunks
 + speed

VHIT - suspends the exec until a child finishes (any, or a specific one)
 EXEC - replace the whole image with the given command
 EXIT - exits the current proc and goes back to the parent

CAUSES OF DL!
 1. Mutual exclusion
 2. Hold and wait
 3. No preemption
 4. Circular wait

SOLUTION TO DL
 - reset → pick a victim proc → create a backup state
 DETECTION OF DL
 graf: □ proc, ○ res.
 → is held by, -- → is req. cycle → deadlock
 PREVENTION OF DL
 → pick an order for and always lock in the order

ZOMBIE → process that has finished execution but still has an entry in the proc table; wait to avoid or before creating: signal (SIGCHLD, SIG-IGN)
 ORPHAN → still execute parent is dead

SCHEDULING ALGO
 ① FIRST COME, FIRST SERVED
 ② SHORTEST FIRST: we need to know durations
 ③ not always right
 ④ Priority Based
 + starvation
 ⑤ Deadline Scheduling based on the finish time (must finish by)

SEMAPHORE (value, queue)
 V(S) - SIGNAL for A:
 V(S)++
 if (V(S) <= 0)
 q(s) → B
 If B comes through
 state(s) = Ready
 pass ctrl to sig.
 scheduler
 else
 pass ctrl to A

Process vs. Thread
 P contains one or more T
 ⊕ quicker to create T > P
 ⊕ quicker to switch between T
 ⊕ T share data easily
 ⊖ no security between T
 ⊖ one T block ⇒ all T in that P block

UNIX - c: counts
 - u: displays
UNIX filesystem structure

Block	Content
0	boot block
1	super block
2	inode blocks
...	...
n	...
n+1	Data blocks

⊗ a kind of map for
FOR I-NODES addr
 - 1-10: for the first 10 blocks of 512 bytes of the file
 - 11: indirectare simple (the next 128 blocks of 512 bytes)
 - 12: indirectare dubla (128² * 512)
 - 13: indirectare tripla (128³ * 512) (B/A)^m

CACHE ① Direct Cache:
 store the page in location
 page-addr % cache-size
 ⊕ fast ⊕ simple
 ⊖ cache trashing (collisions)
 ② Associative Cache
 place each page on a free location found through iterating the cache ⊖ slow

CACHE ③ Set Associative Cache
 - hash table; cache org in sets of pages
 - the set of a page is calc. as page-addr % cache-set, then a free location is found iterating through the set
 ⊕ flexibility ⊕ avoid col.
 ⊖ doesn't use all available cache lines

SIGNAL - acts a function to be called for a certain signal (EXIT, DEATH, etc)
 * does not signal anything

MOUNT associate a storage device to a particular location in the directory tree
 int dup (int dir, Vch)
 int dup (int d, vche, int d-mount)

Read la FIFO: waits until there is no proc open for writing or until there is data; returns the read data or 0 if & a proc to write

int open(char* name, int flag, [int rights])
 O_RDONLY, O_WRONLY, O_RDWR, O_APPEND, O_CREAT, O_TRUNC, O_NODELAY (FIFO)
 read/write (handle, buff, nr-bytes)
 WAIT → RUN?
 → program loaded into mem waiting for CPU (OS scheduler decide)
 RUN → WAIT
 → additional cond.

How many threads to process 1M files? A: the number of available CPUs (cores of the processor)
 Why do I/O ops cause a proc to move from state RUN to WAIT?
 A: The proc. has to pause itself while it gets/sends the data through the I/O ops and it frees the proc in the machine for other procs.

How is the address calculation done in the address fixed part?
 A: multiply prod. part. and the RAM locations are sorted.
 first prog → first position, first add → beginning of part

line containing a null not to
 grp "a" filename | grp v b
 files

A → B → C → A
 if (a == 0) {
 open("x", O_RDONLY);
 open("z", O_WRONLY);
 if (b == 0) ... x, y
 if (c == 0) ... z, y
 2 sets of 4 vs 4 sets of 2
 at least 3 vowels
 grp -E ^(*[aeiou]){3}.*\$

4 rules that specify the no of reps
 + → 1 &
 * → 0 &
 ? → 0/1
 {n} - n times

FIFO for reading; proc?
 ∞ wait until a prog opens the read file.

every open pipe → DON'T forget to close the pipe

Hard-link only in the same partition → when creating a hard link, the OS creates a new directory entry that points to the existing inode of the file; it works similarly in the same partitions i.e. the inode is unique in that part; the OS cannot assign that inode no. to a new diff. part because parts have unique inodes within them.

at least x vowels/numbers.
 grp -E ^^[^']*+([aeiou][^']*+){m}\$
 having a group:
 grp -E ^^[^']*+([aeiou][^']*+){m}\$
 all odd no of char
 grp -E ^^[^']*+([aeiou][^']*+){m}\$
 swap neighbouring digits
 md -E 'sK[0-9][0-9]/[2]/[1]/[9]'
 multiple of binary 4
 grp -E '.*([01]*0523){24}.*'
 binary multiple of 4 100
 5 digits